

Credit Card Default Prediction

Risk, interpretation and decision costs

Emerging Topics in FinTech

THE TEAM

Lei Peng

G2509090C · Part 1

Zhang Hanyu

G2509091L · Part 2

Zhou Xinzhe

G2509033F · Part 3

Xu Yiqun

G2509092H · Part 4

Contribution: 25% each

How we split and use the data

30,000 clients • 23 predictors •
22.12% default

Categories learned within each
training fold

Validation selects policies; historical
test reports outcomes

Partition and purpose

Fitting	18,000
----------------	--------

Five-fold CV and final fitting

Validation	6,000
-------------------	-------

Select decision thresholds

Test	6,000
-------------	-------

Evaluate frozen decisions

Source: data/splits; protocol_frozen.json; results/final/audit.json

Logistic regression: the linear reference

L2 model + fold-trained scaling + one-hot categories

12 declared settings; CV selected
C=0.1, unweighted

Original LR notebook was empty;
original process unverified

AP 0.4967


Baseline AP 0.4970

No ranking improvement from selection
LR* is a separate post-hoc supplement

Source: results/final/model_comparison.csv

Decision tree: controlling overfitting

Test average precision



21 configurations: depth × minimum leaf size

Unrestricted baseline fits training folds almost perfectly

Selected minimum leaf size: 100 observations

A large baseline gap signals overfitting.
CV selection improves this tree's ranking.

Source: results/final/model_comparison.csv

Reading the tree's first decision

Root rule: recent repayment status
 $\text{PAY_0} \leq 1.5$

Selected tree: depth 18, with 134
leaves

Node proportions describe this fitting
sample

Recent repayment status

$\text{PAY_0} \leq 1.5$


Yes · 16.51% No · 70.13%

Default proportion in each fitting-sample branch.
These are associations, not causal effects.

Source: saved decision-tree model; integration audit/report evidence.

Ranking performance across four models

Test average precision



AP emphasizes positive-class retrieval

AP is not trapezoidal PR-AUC

CV selection precedes these test comparisons


Baseline and selected models are both retained.

Small ensemble differences do not establish superiority.

Source: results/final/model_comparison.csv

Random forest: averaging and regularization

Test average precision



Baseline + 23 seeded candidate settings

Selected: 500 trees, depth 8, minimum leaf 2

Positive-class weight 3; feature fraction 0.5

Averaging reduces dependence on one tree.

Weights also change the score distribution.

Source: results/final/model_comparison.csv

Random forest: the review workload

Validation selected $r=5$ threshold:
0.3229

Recall: 52.0% → 77.1%

Alert rate: 21.1% → 48.4%

3,758 → 3,404

Test cost = $FP + 5 \times FN$

Error cases checked under each policy

Same extreme cases can legitimately recur

Source: results/final/model_comparison.csv

What makes the comparison credible

Identical fit / validation / five-fold memberships

Metrics recomputed after row-ID alignment

Frozen thresholds use score $\geq$ threshold at full precision

Source: results/final/audit.json; integration/LOGISTIC_PROVENANCE.md

Shared evidence, different searches


Shared data and objective

Different search budgets

Historical exposure and LR* caveat remain

XGBoost: no gain in test ranking

Test average precision



24 configurations; select by CV AP

Selected: 150 trees, depth 4, rate 0.03

Selected ROC-AUC: 0.7762


Tuned – baseline AP: -0.0028

Conditional 95% interval: [-0.0106, 0.0054]

Source: results/final/model_comparison.csv

XGBoost: fewer misses, more reviews

Test confusion trade-off



Validation r=5 cutoff: 0.3150

Recall 54.9% → 82.5%; alerts 55.7%

Cost 3,685 → 3,408: 7.5% reduction

Same ranking scores, different action policy.

Review costs are hypothetical; capacity is unmeasured.

Source: results/final/model_comparison.csv

What the results mean for credit review

Ranking, explanations, and workload need separate evidence

No clear ensemble winner from these small differences

Next: fresh temporal data, calibration, subgroup errors

Discussion

Set realistic costs and review capacity.
Validate the policy on fresh data.

Historical educational study; no causal or deployment-readiness claim.

Backup A • AP, ROC-AUC and threshold metrics

AP = sum of recall increments × precision

ROC-AUC measures positive/negative score ordering

Precision, recall, F1 and accuracy require a threshold

Two different definitions

Saved continuous scores supply ranking metrics.

Changing the cutoff changes decisions only.

Source: [scikit-learn average_precision_score documentation](#); part4/evaluation.py

Backup B • Cost ratios and fixed-prediction uncertainty

Predeclared ratios: $r = 1, 3, 5, 10$

Validation minimum cost; ties prefer largest cutoff

1,000 stratified paired bootstrap resamples

What the intervals cover

Fixed models / thresholds / class counts

Excludes fitting / selection uncertainty

Excludes population / prevalence shifts

Source: frozen protocols; integration/uncertainty.py; bootstrap_intervals.csv

Backup C • Data provenance and evidence boundaries

UCI cleaning reproduced exactly from downloaded source

EDUCATION 0/5/6 → 4; MARRIAGE
0 → 3

Retained duplicate vectors can span partitions

Sources: results/final/source_audit.json; integration/LOGISTIC_PROVENANCE.md

The limits of historical evidence

Original LR process remains unverified.

LR* source was committed before its one run.

No new unseen-test claim is made.

Backup D • Sources, responsibilities and reproducibility

UCI: doi.org/10.24432/C55S3H

scikit-learn: [average_precision_score](https://scikit-learn.org/stable/modules/average_precision_score.html)
documentation

XGBoost paper:
doi.org/10.1145/2939672.2939785

Member contributions • 25% each

Lei Peng • G2509090C

Part 1 — Data and logistic regression

Zhang Hanyu • G2509091L

Part 2 — Decision tree and model comparison

Zhou Xinzhe • G2509033F

Part 3 — Random forest and experiment audit

Xu Yiqun • G2509092H

Part 4 — XGBoost and cost analysis

Sources: Group 5 repository, frozen protocols, row-level predictions and integration audit.